


Introduction to Probability and Statistics



Study Note

Han Cheol Moon
School of Computer Science and Engineering
Nanyang Technological University
Singapore
hancheol001@e.ntu.edu.sg
June 24, 2024

Contents

1	Introduction to Probability	1
1.1	Introduction	1
2	Combinatorics	2
2.1	Multiplication Principle	2
2.2	Ordered Sampling with Replacement	2
2.3	Ordered Sampling without Replacement: Permutations	2
2.4	Unordered Sampling without Replacement: Combinations	3
2.5	Bernoulli Trials and Binomial Distribution	4
2.6	Unordered Sampling with Replacement	5
3	Discrete Random Variables	7
3.1	Random Variables	7
3.2	Probability Mass Function (PMF)	7
3.3	Special Distributions	7
3.3.1	Bernoulli Distribution	7
3.3.2	Geometric Distribution	8
3.3.3	Binomial Distribution	8
3.3.4	Hyper-Geometric Distribution	10
3.3.5	Poisson Distribution	10
3.3.6	Poisson as an Approximation for Binomial	11
3.4	Cumulative Distribution Function	11
3.5	Expectation	12

3.6 Functions of Random Variables	12
3.6.1 Expected Value of a Function of a Random Variable (LOTUS)	13
3.7 Variance	13
3.8 Standard Deviation	13
4 Continuous and Mixed Random Variables	14
4.1 Introduction	14
4.2 Probability Density Function (PDF)	14
4.3 Expected Value and Variance	15
4.3.1 Expected Value of a Function of a Continuous Random Variable	16
4.3.2 Variance	16
4.4 Functions of Continuous Random Variables	16
4.5 The Method of Transformations	17
4.5.1 Intuitive Explanation	18

Chapter 1

Introduction to Probability

1.1 Introduction

The probability of event A as

$$P(A) = \frac{\text{Number of times A occurs}}{\text{Total number of outcomes}}$$

This commonsense understanding of probability is called the *relative frequency definition*.

Chapter 2

Combinatorics

2.1 Multiplication Principle

Suppose that we perform r experiments such that the k -th experiment has n_k possible outcomes, for $k = 1, 2, \dots, r$. Then there are a total of $n_1 \times n_2 \times n_3 \times \dots \times n_r$ possible outcomes for the sequence of r experiments.

2.2 Ordered Sampling with Replacement

Here we have a set with n elements (*e.g.*, $A = \{1, 2, 3, \dots, n\}$), and we want to draw k samples from the set such that ordering matters and repetition is allowed. For example, if $A = \{1, 2, 3\}$ and $k = 2$, there are 9 different possibilities. In general, we can argue that there are k positions in the chosen list: (Position 1, Position 2, \dots , Position k). There are n options for each position. Thus, when ordering matters and repetition is allowed, the total number of ways to choose k objects from a set with n elements is

$$n \times n \times \dots \times n = n^k.$$

2.3 Ordered Sampling without Replacement: Permutations

Consider the same setting as above, but now repetition is not allowed. For example, if $A = \{1, 2, 3\}$ and $k = 2$, there are 6 different possibilities. In general, we can argue that there are k positions in the chosen list: (Position 1, Position 2, \dots , Position k). There are n options for the first position, $(n - 1)$ options for the second position (since one element has already been allocated to the first position and cannot be chosen here), $(n - 2)$ options for the third position, and $(n - k + 1)$ options for the k -th position. Thus, when ordering matters and repetition is not allowed, the total number of ways to choose k objects from a set with n elements is

$$n \times (n - 1) \times \dots \times (n - k + 1).$$

It is called a k permutation of the elements in set A . We use the following notation:

$$P_k^n = n \times (n - 1) \times \dots \times (n - k + 1).$$

Note that if k is larger than n , then $P_k^n = 0$.

Example: Birthday problem or birthday paradox is a problem that If k people are at a party, what is the probability that at least two of them have the same birthday? Suppose that there are $n = 365$ days in a year and all days are equally likely to be the birthday of a specific person.

$$P(A) = 1 - \frac{P_k^n}{n^k}.$$

The reason this is called a paradox is that $P(A)$ is numerically different from what most people expect. For example, if there are $k = 23$ people in the party, what do you guess is the probability that at least two of them have the same birthday, $P(A)$? The answer is .5073, which is much higher than what most people guess. The probability crosses 99 percent when the number of peoples reaches 57. But why is the probability higher than what we expect?

It is important to note that in the birthday problem, neither of the two people are chosen beforehand. To better answer this question, let us look at a different problem: I am in a party with $k - 1$ people. What is the probability that at least one person in the party has the same birthday as mine? Well, we need to choose the birthdays of $k - 1$ people, the total number of ways to do this is n^{k-1} . The total number of ways to choose the birthdays so that no one has my birthday is $(n - 1)^{k-1}$. Thus, the probability that at least one person has the same birthday as mine is

$$P(B) = 1 - \left(\frac{n-1}{n} \right)^{k-1}.$$

Now, if $k = 23$, this probability is only $P(B) = 0.0586$, which is much smaller than the corresponding $P(A) = 0.5073$. The reason is that event B is looking only at the case where one person in the party has the same birthday as me. This is a much smaller event than event A which looks at all possible pairs of people. Thus, $P(A)$ is much larger than $P(B)$. We might guess that the value of $P(A)$ is much lower than it actually is, because we might confuse it with $P(B)$.

Permutations of n elements: An n -permutation of n elements is just called a permutation of those elements. In this case $k = n$ and we have

$$\begin{aligned} P_n^n &= n \times (n-1) \times \cdots \times (n-n+1) \\ &= n \times (n-1) \times \cdots \times 1, \end{aligned}$$

which is denoted $n!$. We can rewrite as

$$P_k^n = \frac{n!}{(n-k)!}.$$

2.4 Unordered Sampling without Replacement: Combinations

Here we have a set with n elements, e.g., $A = \{1, 2, 3, \dots, n\}$ and we want to draw k samples from the set such that ordering does not matter and repetition is not allowed. Thus, we basically want to choose a k -element subset of A , which we also call a k -combination of the set A . For example if $A = \{1, 2, 3\}$ and $k = 2$, there are 3 different possibilities. We show the number of k -element subsets of A by

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}.$$

This is also called the *binomial coefficient*. This is because the coefficients in the binomial theorem are given by

$$(a + b)^n = \sum_{k=0}^n \binom{n}{k} a^k b^{n-k}.$$

An intuitive way to understand this is that there are $n \times (n - 1) \times \cdots \times (n - k + 1)$ ways to place items and the $k \times \cdots \times 1$ ways to order the times, which can be ignored.

A simple way to find $\binom{n}{k}$ is to compare it with P_k^n . Note that the difference between the two is ordering.

$$P_k^n = \binom{n}{k} \times k!.$$

Example 1: I choose 3 cards from the standard deck of cards. What is the probability that these cards contain at least one ace?

- The sample space contains all possible ways to choose 3 cards from 52 cards.
- There are $52 - 4 = 48$ non-ace cards

Example 2: How many distinct sequences can we make using 3 letter “A”s and 5 letter “B”s? (AAABBBBB, AABBBBB, .)

You can think of this problem in the following way. You have $3+5=8$ positions to fill with letters A or B. From these 8 positions, you need to choose 3 of them for “A”s. Whatever is left will be filled with “B”s. Thus the total number of ways is

$$\binom{8}{3}.$$

Equivalently, you should have chosen the locations for Bs.

$$\binom{8}{5}.$$

The same argument can be repeated for general n and k to conclude

$$\binom{n}{k} = \binom{n}{n - k}.$$

2.5 Bernoulli Trials and Binomial Distribution

A *Bernoulli Trial* is a random experiment that has two possible outcomes which we can label as “success” and “failure”, such as

- You toss a coin. The possible outcomes are H and T.

We usually denote the probability of success by p and probability of failure by $q = 1 - p$. If we have an experiment in which we perform n independent Bernoulli trials and count the total number of successes, we call it a binomial experiment. For example, you may toss a coin n times repeatedly and be interested in the total number of heads.

Example: Suppose that I have a coin for which $P(H) = p$ and $P(T) = 1 - p$. I toss the coin 5 times.

- $P(THHHH) = p(T) \times p(H) \cdots = (1 - p)p^4$
- $P(HTHHH) = (1 - p)p^4$
- $P(HHTHH) = (1 - p)p^4$
- $B = \{THHHH, HTHHH, HHTHH, HHHTH, HHHHT\}, P(B) = 5p^4(1 - p)$
- Let $C = \{TTHHH, THTHH, \dots\}$.

$$\begin{aligned} P(C) &= P(TTHTH) + P(THTHH) + \dots \\ &= |C|p^3(1 - p)^2 \end{aligned}$$

- The $|C|$ is the total number of distinct sequences that you can create using two tails and three heads.

$$\binom{5}{3}.$$

- Therefore,

$$P(C) = \binom{5}{3}p^3(1 - p)^2$$

Now we can define *Binomial Formula*: For n independent Bernoulli trials where each trial has success probability p , the probability of k successes is given by

$$P(k) = \binom{n}{k}p^k(1 - p)^{n-k}.$$

Similarly, *multinomial coefficients* is given by

$$\binom{n}{n_1, n_2, \dots, n_r} = \frac{n!}{n_1!n_2!\dots n_r!}.$$

2.6 Unordered Sampling with Replacement

Suppose that we want to sample from the set $A = \{a_1, a_2, \dots, a_n\}$ k times such that repetition is allowed and ordering does not matter. For example, if $A = \{1, 2, 3\}$ and $k = 2$, then there are 6 different ways of doing this.

How can we get the number 6 without actually listing all the possibilities? One way to think about this is to note that any of the pairs in the above list can be represented by the number of 1's, 2's and 3's it contains. That is, if x_1 is the number of ones, x_2 is the number of twos, and x_3 is the number of threes, we can equivalently represent each pair by a vector (x_1, x_2, x_3) , i.e.,

- $(1, 1) \rightarrow (x_1, x_2, x_3) = (2, 0, 0)$
- $(1, 2) \rightarrow (x_1, x_2, x_3) = (1, 1, 0)$
- $(2, 3) \rightarrow (x_1, x_2, x_3) = (0, 1, 1)$

Note that here $x_i \geq 0$ are integers and $x_1 + x_2 + x_3 = 2$. Thus, we can claim that the number of ways we can sample two elements from the set $A = \{1, 2, 3\}$ such that ordering does not matter and repetition is allowed is the same as solutions to the following equation

$$x_1 + x_2 + x_3 = 2,$$

where $x_i \in \{0, 1, 2\}$. We can generalize this by saying: The total number of distinct k samples from an n -element set such that repetition is allowed and ordering does not matter is the same as the number of distinct solutions to the equation

$$x_1 + x_2 + \cdots + x_n = k,$$

where $x_i \in \{0, 1, 2, \dots\}$. The number of distinct solution to the equation is given by

$$\binom{n+k-1}{k} = \binom{n+k-1}{n-1}.$$

Proof 1 Let us first define following simple mapping in which we replace an integer x_i with vertical lines i.e., |. For instance, $x_1 + x_2 + x_3 = 2$, then we can equivalently write |++| for 1+0+1. We have an unique representation using vertical lines and plus signs. Each solution can be represented by k vertical lines and $n-1$ plus sings. Thus, we get

$$\binom{n-1+k}{k} = \binom{n-1+k}{n-1}.$$

Chapter 3

Discrete Random Variables

3.1 Random Variables

A random variable X is a function from the sample space to the real numbers.

$$X : S \rightarrow \mathbb{R}$$

3.2 Probability Mass Function (PMF)

If X is a discrete random variable then its range R_X is a countable set, so, we can list the elements in R_X . In other words, we can write

$$R_X = \{x_1, x_2, \dots\}$$

Note that here x_1, x_2, \dots are possible values of the random variable X . While random variables are usually denoted by capital letters, to represent the numbers in the range we usually use lowercase letters. For a discrete random variable X , we are interested in knowing the probabilities of $X = x_k$.

Let X be a discrete random variable with range $R_X = \{x_1, x_2, \dots\}$ (finite or countably infinite). The function

$$P_X(x_k) = P(X = x_k), \text{ for } k = 1, 2, 3 \dots$$

is called the probability mass function (PMF) of X .

3.3 Special Distributions

3.3.1 Bernoulli Distribution

$$P_X(x) = \begin{cases} p & \text{for } x = 1 \\ 1 - p & \text{for } x = 0 \\ 0 & \text{Otherwise} \end{cases}$$

A Bernoulli random variable is associated with a certain event A . If event A occurs (for example, if you pass the test), then $X = 1$; otherwise $X = 0$. For this reason the Bernoulli random variable, is also called the *indicator* random variable.

3.3.2 Geometric Distribution

Suppose that I have a coin with $P(H) = p$. I toss the coin until I observe the first heads. We define X as the total number of coin tosses in this experiment. Then X is said to have geometric distribution with parameter p . In other words, you can think of this experiment as **repeating independent Bernoulli trials until observing the first success**. The range of X here is $R_X = \{1, 2, 3, \dots\}$.

$$P_X(k) = P(X = k) = (1 - p)^{k-1}p, \text{ for } k = 1, 2, 3, \dots$$

3.3.3 Binomial Distribution

Suppose that I have a coin with $P(H) = p$. I toss the coin n times and define X to be the total number of heads that I observe. Then X is binomial with parameter n and p . The range of X in this case is $R_X = 0, 1, 2, \dots, n$.

$$P_X(k) = P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k}.$$

Here is a useful way of thinking about a binomial random variable. It can be obtained by n independent coin tosses. If we think of each coin toss as a Bernoulli random variable, the $\text{Binomial}(n, p)$ random variable is a sum of n independent $\text{Bernoulli}(p)$ random variables. This is stated more precisely in the following lemma.

If X_1, X_2, \dots, X_n are independent $\text{Bernoulli}(p)$ random variables, then the random variable X defined by $X = X_1 + X_2 + \dots + X_n$ has a $\text{Binomial}(n, p)$ distribution.

Example:

- Let $X \sim \text{Binomial}(n, p)$ and $Y \sim \text{Binomial}(m, p)$ be two independent random variables. Define a new random variable as $Z = X + Y$. Find the PMF of Z .
- Solution 1: Since $X \sim \text{Binomial}(n, p)$, we can think of X as the number of heads in n independent coin tosses:

$$X = X_1 + \dots + X_n,$$

where the X_i 's are independent Bernoulli RVs. Similarly, $Y \sim \text{Binomial}(m, p)$. Thus, the RV $Z = X + Y$ will be the total number of heads in $n + m$ coin tosses:

$$Z = X + Y = X_1 + \dots + X_n + Y_1 + \dots + Y_m.$$

Therefore, Z is a binomial RV with parameters $m + n$ and p , i.e., $\text{Binomial}(m + n, p)$.

- Solution 2: First, we note that $R_z = \{0, 1, \dots, m + n\}$. For $k \in R_z$, we get

$$P_Z(k) = P(Z = k) = P(X + Y = k).$$

We will find $P(X + Y = k)$ by using conditioning and the law of total probability.

$$\begin{aligned}
P(Z = k) &= P(X + Y = k) \\
&= \sum_{i=0}^n P(X + Y = k | X = i) P(X = i) \\
&= \sum_{i=0}^n P(Y = k - i | X = i) P(X = i) \\
&= \sum_{i=0}^n P(Y = k - i) P(X = i) \quad \text{Since } X \text{ and } Y \text{ are independent} \\
&= \sum_{i=0}^n \binom{m}{k-i} p^{k-i} (1-p)^{m-k+i} \binom{n}{i} p^i (1-p)^{n-i} \\
&= \sum_{i=0}^n \binom{m}{k-i} \binom{n}{i} p^k (1-p)^{m+n-k} \\
&= p^k (1-p)^{m+n-k} \sum_{i=0}^n \binom{m}{k-i} \binom{n}{i} \\
&= \binom{m+n}{k} p^k (1-p)^{m+n-k} \quad \text{by Vandermonde's identity}
\end{aligned}$$

Negative Binomial (Pascal) Distribution The *negative binomial* or *Pascal distribution* is a **generalization of the geometric distribution**. It relates to the random experiment of **repeated independent trials until observing m successes**. Suppose that I have a coin with $P(H) = p$. I toss the coin until I observe m heads, where $m \in \mathbb{N}$. We define X as the total number of coin tosses in this experiment. Then X is said to have Pascal distribution with parameter m and p . We write $X \sim \text{Pascal}(m, p)$. Note that $\text{Pascal}(1, p = \text{Geometric}(p))$, since the geometric distribution repeats trials until observing the first success. Note that by our definition the range of X is given by $R_X = \{m, m+1, m+2, m+3, \dots\}$, since X is the number of coin tosses to observe m target events.

Let's derive the PMF of a $\text{Pascal}(m, p)$ RV X . To find the probability of the event $A = \{X = k\}$, we argue as follows. By definition, event A can be written as $A = B \cap C$, where

- B is the event that we observe $m - 1$ heads (*i.e.*, successes) in the first $k - 1$ trials
- C is the event that we observe a head in the k -th trial.

Note that B and C are independent events because they are related to different independent trials (coin tosses). Thus,

$$P(A) = P(B \cap C) = P(B)P(C).$$

We get $P(C) = p$, so

$$P(B) = \binom{k-1}{m-1} p^{m-1} (1-p)^{(k-1)-(m-1)} = \binom{k-1}{m-1} p^{m-1} (1-p)^{k-m}.$$

Finally, we obtain

$$P(A) = \binom{k-1}{m-1} p^m (1-p)^{k-m}.$$

3.3.4 Hyper-Geometric Distribution

You have a bag that contains b blue marbles and r red marbles. You choose $k \leq b + r$ marbles at random (without replacement). Let X be the number of blue marbles in your sample. By this definition, we have $X \leq \min(k, b)$. Also, the number of red marbles in your sample must be less than or equal to r , so we conclude $X \geq \max(0, k - r)$. Therefore, the range of X is given by $R_X = \{\max(0, k - r), \max(0, k - r) + 1, \max(0, k - r) + 2, \dots, \min(k, b)\}$.

To find $P_X(x)$, note that total number of ways to choose k marbles from $b + r$ marbles is $\binom{b+r}{k}$. The total number of ways to choose x blue marbles and $k - x$ red marbles is $\binom{b}{x} \binom{r}{k-x}$. Thus, we get

$$P_X(x) = \frac{\binom{b}{x} \binom{r}{k-x}}{\binom{b+r}{k}}, \quad \text{for } x \in R_X.$$

3.3.5 Poisson Distribution

The Poisson distribution is one of the most widely used probability distributions. It is usually used in scenarios where we are **counting the occurrences of certain events in an interval of time or space**. In practice, it is often an approximation of a real-life random variable. Here is an example of a scenario where a Poisson random variable might be used. Suppose that we are counting the number of customers who visit a certain store from 1pm to 2pm. Based on data from previous days, we know that on average $\lambda = 15$ customers visit the store. Of course, there will be more customers some days and fewer on others. Here, we may model the random variable X showing the number of customers as a Poisson random variable with parameter $\lambda = 15$. Let us introduce the Poisson PMF first, and then we will talk about more examples and interpretations of this distribution.

$$P_X(k) = e^{-\lambda} \frac{\lambda^k}{k!}.$$

Note that λ is the mean number of events within a given interval of time or space.

Example: The number of emails that I get in a weekday can be modeled by a Poisson distribution with an average of 0.2 emails per minute.

- What is the probability that I get no emails in an interval of length 5 minutes?
 - For 5 minutes, there would be 1 email on average. Thus, $\lambda = 1$,
- $$P(X = 0) = P_X(0) = e^{-\lambda} \frac{\lambda^0}{0!} = \frac{1}{e} \approx 0.37$$
- What is the probability that I get more than 3 emails in an interval of length 10 minutes?
 - Let Y be the number of emails that I get in the 10-minute interval. Then by the assumption Y is a Poisson RV with parameter $\lambda = 10 \times 0.2 = 2$. Thus,

$$\begin{aligned} P(Y > 3) &= 1 - P(Y \leq 3) \\ &= 1 - (P_Y(0) + P_Y(1) + P_Y(2) + P_Y(3)) \\ &= 1 - e^{-\lambda} - \frac{e^{-\lambda}\lambda}{1!} - \frac{e^{-\lambda}\lambda^2}{2!} - \frac{e^{-\lambda}\lambda^3}{3!} \\ &\approx 0.1429 \end{aligned}$$

Imagine you have a busy customer service center that receives phone calls. You want to know how many calls to expect in an hour, but calls can come at any moment and don't follow a strict schedule.

- Average Rate (λ): First, you determine the average number of calls you receive per hour. Let's say it's 10 calls per hour. This average rate is denoted by the symbol λ
- Probability Calculation: Using the Poisson formula, you can calculate the probability of receiving a certain number of calls in any given hour.

$$P(X = k) = \frac{e^{-\lambda} \lambda^k}{k!}$$

- $P(X = k)$ is the probability of getting k calls in an hour.
- e is the base of the natural logarithm (approximately equal to 2.71828).
- λ is the average rate (10 calls per hour).
- k is the number of calls you want to find the probability for.
- $k!$ (k factorial) is the product of all positive integers up to k .

3.3.6 Poisson as an Approximation for Binomial

The Poisson distribution can be viewed as the limit of binomial distribution. Suppose $X \sim \text{Binomial}(n, p)$ where n is very large and p is very small. In particular, assume that $\lambda = np$ is a positive constant. We show that the PMF of X can be approximated by the PMF of a $\text{Poisson}(\lambda)$ random variable. The importance of this is that Poisson PMF is much easier to compute than the binomial. Let us state this as a theorem.

Let $X \sim \text{Binomial}(n, p = \frac{\lambda}{n})$, where $\lambda > 0$ is fixed. Then for any $k \in \{0, 1, 2, \dots\}$ we have

$$\lim_{n \rightarrow \infty} P_X(k) = \frac{e^{-\lambda} \lambda^k}{k!}.$$

References Poisson

3.4 Cumulative Distribution Function

The PMF is one way to describe the distribution of a discrete RV. As we will see later on, PMF cannot be defined for continuous random variables. The cumulative distribution function (CDF) of a random variable is another method to describe the distribution of random variables. The advantage of the CDF is that it can be defined for any kind of RV (discrete, continuous, and mixed).

Definition 1 *Cumulative Distribution Function* The cumulative distribution function (CDF) of random variable X is defined as

$$F_X(x) = P(X \leq x), \forall x \in \mathbb{R}.$$

Note that the subscript X indicates that this is the CDF of the random variable X . Also, note that the CDF is defined for all $x \in \mathbb{R}$.

3.5 Expectation

If you have a collection of numbers a_1, a_2, \dots, a_N , their average is a single number that describes the whole collection. Now, consider a random variable X . We would like to define its average, or as it is called in probability, its expected value or mean. The expected value is defined as the weighted average of the values in the range.

Definition 2 *Expected Value* Let X be a discrete RV with range $R_X = \{x_1, x_2, \dots\}$. The expected value of X , denoted by EX is defined as

$$EX = \sum_{x_k \in R_X} x_k P(X = x_k) = \sum_{x \in R_X} x_k P_X(x_k)$$

3.6 Functions of Random Variables

If X is a RV and $Y = g(X)$, then Y itself is a random variable. Thus, we can talk about its PMF, CDF, and expected value. First note that the range of Y can be written as

$$R_Y = \{g(x) | x \in R_X\}.$$

If we already know the PMF of X , to find the PMF of $Y = g(X)$, we can write

$$\begin{aligned} P_Y(y) &= P(Y = y) \\ &= P(g(X) = y) \\ &= \sum_{x:g(x)=y} P_X(x) \end{aligned}$$

Example: Let X be a discrete RV with $P_X(k) = \frac{1}{5}$ for $k = -1, 0, 1, 2, 3$. Let $Y = 2|X|$. Find the range and PMF of Y . The range of Y is

$$\begin{aligned} R_Y &= \{2|x|\} \\ &= \{0, 2, 4, 6\} \end{aligned}$$

To find $P_Y(y)$, we need to find $P(Y = y)$ for $y = 0, 2, 4, 6$:

$$\begin{aligned} P_Y(0) &= P(Y = 0) = P(2|x| = 0) \\ &= P(X = 0) = \frac{1}{5} \\ P_Y(2) &= P(Y = 2) = P(2|x| = 2) \\ &= P(X = -1 \text{ or } X = 1) \\ &= P_X(-1) + P_X(1) = \frac{2}{5} \\ &\vdots \end{aligned}$$

3.6.1 Expected Value of a Function of a Random Variable (LOTUS)

Let X be a discrete random variable with PMF $P_X(x)$, and let $Y = g(X)$. Suppose that we are interested in finding EY . One way to find EY is to first find the PMF of Y and then use the expectation formula $EY = E[g(X)] = \sum_{y \in R_y} y P_Y(y)$. But there is another way which is usually easier. It is called the *law of the unconscious statistician* (LOTUS).

$$\mathbb{E}[g(X)] = \sum_{x_k \in R_X} g(x_k) P_X(x_k)$$

One of the main points of the theorem is that you can compute $\mathbb{E}[g(X)]$ without computing $P_Y(y)$. In practice it is usually easier to use LOTUS than direct definition when we need $\mathbb{E}[g(X)]$.

3.7 Variance

The variance of a random variable X , with mean $EX = \mu_X$, is defined as

$$\text{Var}(X) = \mathbb{E}[(X - \mu_X)^2].$$

To compute $\text{Var}(X) = \mathbb{E}[(X - \mu_X)^2]$, note that we need to find the expected value of $g(X) = (X - \mu_X)^2$, so we can use **LOTUS**. In particular, we can write

$$\text{Var}(X) = \mathbb{E}[(X - \mu_X)^2] = \sum_{x_k \in R_X} (x_k - \mu_X)^2 P_X(x_k).$$

3.8 Standard Deviation

$$\text{SD}(X) = \sigma_X = \sqrt{\text{Var}(X)}.$$

A useful formula for computing the variance is

$$\text{Var}(X) = \mathbb{E}[X^2] - \mathbb{E}[X]^2.$$

We can find $\mathbb{E}[X^2]$ using LOTUS:

$$\mathbb{E}[X^2] = \sum_{x_k \in R_X} x_k^2 P_X(x_k).$$

Chapter 4

Continuous and Mixed Random Variables

4.1 Introduction

Remember that discrete random variables can take only a countable number of possible values. On the other hand, a continuous random variable X has a range in the form of an interval or a union of non-overlapping intervals on the real line (possibly the whole real line). Also, for any $x \in \mathbb{R}$, $P(X = x) = 0$. Thus, we need to develop new tools to deal with continuous random variables. The good news is that the theory of continuous random variables is completely analogous to the theory of discrete random variables. Indeed, if we want to oversimplify things, we might say the following: take any formula about discrete random variables, and then replace sums with integrals, and replace PMFs with probability density functions (PDFs), and you will get the corresponding formula for continuous random variables.

4.2 Probability Density Function (PDF)

To determine the distribution of a discrete random variable we can either provide its PMF or CDF. For continuous random variables, the CDF is well-defined so we can provide the CDF. However, **the PMF does not work for continuous random variables, because for a continuous random variable $P(X = x) = 0$, $\forall x \in \mathbb{R}$.** Instead, we can usually define the probability density function (PDF). The PDF is the *density* of probability rather than the probability mass. The concept is very similar to mass density in physics: its unit is probability per unit length. To get a feeling for PDF, consider a continuous random variable X and define the function $f_X(x)$ as follows (wherever the limit exists):

$$f_X(x) = \lim_{\Delta \rightarrow 0^+} \frac{P(x < X \leq x + \Delta)}{\Delta}.$$

The function $f_X(x)$ gives us the probability density at point x . It is the limit of the probability of the interval $(x, x + \Delta]$ divided by the length of the interval as the length of the interval goes to 0. Remember that

$$P(x < X \leq x + \Delta) = F_X(x + \Delta) - F_X(x).$$

Thus, we get

$$\begin{aligned} f_X(x) &= \lim_{\Delta \rightarrow 0} \frac{F_X(x + \Delta) - F_X(x)}{\Delta} \\ &= \frac{dF_X(x)}{dx} \\ &= F'_X(x), \quad \text{if } F_X(x) \text{ is differentiable at } x. \end{aligned}$$

Let's find the PDF of the uniform random variable $X \sim Uniform(a, b)$, which can be expressed as follows:

$$f_X(x) = \begin{cases} \frac{1}{b-a} & a < x < b \\ 0 & x < a \text{ or } x > b \end{cases}$$

Note that the CDF is not differentiable at points a and b . Nevertheless, this is not important at this moment.

The uniform distribution is the simplest continuous random variable you can imagine. For other types of continuous random variables the PDF is non-uniform. Note that for small values of δ we can write

$$P(x < X \leq x + \delta) \approx f_X(x)\delta.$$

Thus, if $f_X(x_1) < f_X(x_2)$, we can say $P(x_1 < X \leq x_1 + \delta) < P(x_2 < X \leq x_2 + \delta)$, i.e., the value of X is more likely to be around x_2 than x_1 .

Since the PDF is the derivative of the CDF, the CDF can be obtained from PDF by integrations (by assuming absolute continuity):

$$F_X(x) = \int_{-\infty}^x f_X(u)du.$$

Also, we have

$$P(a < X \leq b) = F_X(b) - F_X(a) = \int_a^b f_X(u)du.$$

More generally, for a set A , $P(X \in A) = \int_a^b f_X(u)du$. Note that if we integrate over the entire real line, we must get 1, i.e.,

$$\int_{-\infty}^{\infty} f_X(u)du = 1.$$

4.3 Expected Value and Variance

As we mentioned earlier, the theory of continuous random variables is very similar to the theory of discrete random variables. In particular, usually summations are replaced by integrals and PMFs are replaced by PDFs. The proofs and ideas are very analogous to the discrete case, so sometimes we state the results without mathematical derivations for the purpose of brevity.

Recall that the expected value of a discrete random variable can be obtained as

$$EX = \sum_{x_k \in R_X} x_k P_X(x_k).$$

The expected value of a continuous RV as

$$EX = \int_{-\infty}^{\infty} x f_X(x)dx.$$

4.3.1 Expected Value of a Function of a Continuous Random Variable

Law of the unconscious statistician (LOTUS) for continuous random variables:

$$\mathbb{E}[g(X)] = \int_{-\infty}^{\infty} g(x)f_X(x)dx$$

4.3.2 Variance

$$\begin{aligned}\text{Var}(X) &= \mathbb{E}[(X - \mu_X)^2] \\ &= \int_{-\infty}^{\infty} (x - \mu_X)^2 f_X(x)dx \\ &= EX^2 - (EX)^2 \\ &= \int_{-\infty}^{\infty} x^2 f_X(x)dx - \mu_X^2\end{aligned}$$

Note that for $a, b \in \mathbb{R}$, we always have

$$\text{Var}(aX + b) = a^2\text{Var}(X)$$

4.4 Functions of Continuous Random Variables

If X is a continuous random variable and $Y = g(X)$ is a function of X , then Y itself is a random variable. Thus, we should be able to find the CDF and PDF of Y . It is usually more straightforward to start from the CDF and then to find the PDF by taking the derivative of the CDF. Note that before differentiating the CDF, we should check that the CDF is continuous. As we will see later, the function of a continuous random variable might be a non-continuous random variable. Let's look at an example.

Example: Let X be a $\text{Uniform}(0, 1)$ random variable, and let $Y = e^X$.

- CDF of Y
- PDF of Y
- EY

The PDF of X is given by

$$F_X(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } 0 \leq x \leq 1 \\ 1 & \text{for } x > 1 \end{cases}$$

The range of x , $R_X = [0, 1]$, so the range of Y , $R_Y = [1, e]$. We can find the CDF of Y as follows:

$$\begin{aligned} F_Y(y) &= P(Y \leq y) \\ &= P(e^X \leq y) \\ &= P(X \leq \ln y) \quad , \text{ since } e^x \text{ is an increasing function.} \\ &= F_X(\ln y) \quad \text{by definition.} \\ &= \ln y \quad , \text{ since } F_X(x) = x \text{ for } 0 \leq x \leq 1 \text{ and } 0 \leq \ln y \leq 1. \end{aligned}$$

4.5 The Method of Transformations

So far, we have discussed how we can find the distribution of a function of a continuous random variable starting from finding the CDF. If we are interested in finding the PDF of $Y = g(X)$, and the function g satisfies following properties, it might be easier to use a method called the method of transformations.

- $g(x)$ is differentiable;
- $g(x)$ is a strictly increasing function, that is, if $x_1 < x_2$, then $g(x_1) < g(x_2)$.

Now, let X be a continuous random variable and $Y = g(X)$. We will show that you can directly find the PDF of Y using the following formula.

$$f_Y(y) = \begin{cases} \frac{f_X(x_1)}{g'(x_1)} & \text{where } g(x_1) = y \\ 0 & \text{if } g(x) = y \text{ does not have a solution} \end{cases}$$

Note that since g is strictly increasing, its inverse function g^{-1} is well defined. That is, for each $y \in R_Y$, there exists a **unique** x_1 such that $g(x_1) = y$. We can write $x_1 = g^{-1}(y)$.

$$\begin{aligned} F_Y(y) &= P(Y \leq y) \\ &= P(g(X) \leq y) \\ &= P(X < g^{-1}(y)) \quad \text{since } g \text{ is strictly increasing.} \\ &= F_X(g^{-1}(y)). \end{aligned}$$

To find the PDF of Y , we differentiate,

$$\begin{aligned} f_Y(y) &= \frac{d}{dy} F_X(x_1) \quad \text{by } g(x_1) = y \\ &= \frac{dx_1}{dy} \cdot \underbrace{\frac{d}{dx_1} F_X(x_1)}_{=F'_X(x_1)} \\ &= \frac{dx_1}{dy} f_X(x_1) \\ &= \frac{f_X(x_1)}{g'(x_1)} \quad , \text{since } \frac{dx}{dy} = \frac{1}{\frac{dy}{dx}}. \end{aligned}$$

We can repeat the same argument for the case where g is **strictly decreasing**. In that case, $g'(x_1)$ will be **negative**, so we need to use $|g(x_1)|$. Thus, we can state the following theorem for a *strictly monotonic function*. (A function $g : R \rightarrow R$ is called strictly monotonic if it is strictly increasing or strictly decreasing.)

Bibliography

- [1] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *3rd International Conference on Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings*, 2015.
- [2] Rachit Singh. . https://rachitsingh.com/elbo_surgery/, 2017. Online; accessed 29 January 2014.
- [3] Hany Hassan, Anthony Aue, Chang Chen, Vishal Chowdhary, Jonathan Clark, Christian Federmann, Xuedong Huang, Marcin Junczys-Dowmunt, William Lewis, Mu Li, Shujie Liu, Tie-Yan Liu, Renqian Luo, Arul Menezes, Tao Qin, Frank Seide, Xu Tan, Fei Tian, Lijun Wu, Shuangzhi Wu, Yingce Xia, Dongdong Zhang, Zhirui Zhang, and Ming Zhou. Achieving human parity on automatic chinese to english news translation. *CoRR*, abs/1803.05567, 2018.